

SHUYU GAN

+1 651-410-3956 | gan00067@umn.edu | Minneapolis, MN

shuyugan.github.io | Google Scholar | github.com/shuyugan

EDUCATION

University of Minnesota, Twin Cities

Sep. 2025 – Present

Ph.D. Student in Computer Science

Advisors: **Dongyeop Kang** and **Mingyi Hong**

Zhejiang University

Sep. 2020 – Jun. 2024

Bachelor of Engineering with **Honors** in Computer Science, Chu Kochen Honors College

GPA: **3.94/4.00**

RESEARCH INTERESTS

Reliable Agent Evaluation and Improvement; Inference-Time Scaling & Reasoning; LLM Post-Training (RL); GUI Agents

INDUSTRY RESEARCH EXPERIENCE

Microsoft, Applied Sciences Group

Jun. 2026 – Sep. 2026

Research Intern

Mentor: Yinheng Li

- Investigated autonomous agents for model-training supervision, with a focus on failure-mode detection and diagnosis.
- Explored feedback-guided training interventions as a foundation for reliable recursive self-improvement (RSI); work in progress.

PREPRINTS

SanSi: A Looped Typed Decision Model for System 1.5 Thinking

Shuyu Gan, Young-Jun Lee, Dongyeop Kang

arXiv:2610.07730, 2026

PUBLICATIONS

Do LLMs Recognize Your Latent Preferences? A Benchmark for Latent Information Discovery in Personalized Interaction

Ioannis Tsaknakis, Bingqing Song, **Shuyu Gan**, Dongyeop Kang, Alfredo Garcia, Gaowen Liu, Charles Fleming, Mingyi Hong

COLM 2026

Scaling Unverifiable Rewards: A Case Study on Visual Insights

Shuyu Gan, James Mooney, Pan Hao, Renxiang Wang, Mingyi Hong, Qianwen Wang, Dongyeop Kang

Findings of ACL 2026

A2P-Vis: An Analyzer-to-Presenter Agentic Pipeline for Visual Insights Generation and Reporting

Shuyu Gan, Renxiang Wang, James Mooney, Dongyeop Kang

IEEE VIS 2025 VISxGenAI Workshop; Honorable Mention

Agent S: An Open Agentic Framework that Uses Computers Like a Human

Saaket Agashe, Jiuzhou Han, **Shuyu Gan**, Jiachen Yang, Ang Li, Xin Eric Wang

ICLR 2025

Cultural Self-Adaptive Multimodal Gesture Generation Based on Multiple Culture Gesture Dataset

Jingyu Wu, Shi Chen, **Shuyu Gan**, Weijun Li, Changyuan Yang, Lingyun Sun

ACM Multimedia 2023

ACADEMIC RESEARCH EXPERIENCE

SanSi: A Looped Typed Decision Model for System 1.5 Thinking

Graduate Research Assistant, University of Minnesota

Sep. 2026 – Present
Advisor: Dongyeop Kang

- Proposed SanSi, which turns a pre-trained looped LM (Ouro-1.4B) into a typed decision model that is read and trained (CE + Brier) after every loop, so one model serves every budget from one to eight loops in a single run.
- Built a suite of 10,027 test decisions from 59 sources; with the same data, recipe and seeds, SanSi reached 72.0% accuracy, 13.5 points above a same-shape non-looped model and 5.3 above a newer non-looped model of its size.
- Showed on two depth-controlled tasks that loops extend the solvable depth beyond training, where a $3\times$ larger single-pass model fails; as the only RL reward (no gold answers), SanSi raised a generator’s F1 by 7.7 points.

Selective Test-Time Scaling for Visual Insight Generation

Graduate Research Assistant, University of Minnesota

Jul. 2025 – Jan. 2026
Advisor: Dongyeop Kang

- Developed Selective Test-Time Scaling, a process-based refinement framework that allocates inference-time computation across pipeline stages and prunes weak branches using stage-specific evaluators.
- Built an end-to-end multi-agent data-analysis pipeline for profiling datasets, generating charts and visual insights, and producing reports.
- Evaluated LLM judges against expert preferences using stratified sampling and pairwise annotations; selected the best-aligned judge (Kendall’s $\tau = 0.55$).
- In the main experiment, increased mean judge-rated report quality from 61.64 to 65.86 under a comparable total model-call budget.

Agent S: A General-Purpose Computer-Use Agent

Research Intern, ERIC Lab, University of California, Santa Cruz

Jun. 2024 – Dec. 2024
Advisor: Xin Eric Wang

- Co-developed Agent S, a multimodal language-model agent combining experience-augmented hierarchical planning, continual memory updates, and an Agent-Computer Interface.
- Built and maintained the codebase; designed the ID-based interface; and integrated online search, tree search, world-model reasoning, and accessibility-tree information for visual grounding.
- Achieved state-of-the-art OSWorld performance at publication: 20.58% success versus 11.21% for the baseline (+9.37 percentage points; 83.6% relative improvement).

Concept-Guided Diffusion Model Training

Research Intern, mlPC Lab, University of California, San Diego

Jul. 2023 – Sep. 2023
Advisor: Zhuowen Tu

- Incorporated end-to-end conceptual guidance into diffusion-model training by jointly training a generative adversarial network with a Diffusion Transformer (DiT).
- Trained and fine-tuned DiT baselines on CelebA-HQ with perceptual and adversarial guidance, reducing FID from 13.77 to 11.63.

Culture-Aware Co-Speech Gesture Generation

Research Assistant, International Design Laboratory, Zhejiang University

Feb. 2023 – May 2023
Advisor: Lingyun Sun

- Created and annotated the Multiple Culture Gesture Dataset, comprising 10 cultures and 10,000 motion segments.
- Co-developed a cultural self-adaptive gesture-generation network, implementing text and audio encoders and a cultural feature extractor and introducing adversarial training.
- Improved average FGD from 53.7 to 48.0 and Culture Deception Rate from 33.63% to 39.87%.

TEACHING AND SERVICE

Lead Teaching Assistant, CSCI 5541: Natural Language Processing, University of Minnesota

Fall 2025

Reviewer: NeurIPS 2026 (ED Track); ICML 2026; ACL 2026 Workshop

2026

AWARDS AND HONORS

Outstanding Graduate of Zhejiang University

2024

Honorary Certificate, Chu Kochen Honors College, Zhejiang University

2024

First-Prize Scholarship and Outstanding Student, Zhejiang University (Top 2%)

2022–2023

Top-Notch Scholarship for Basic Disciplines, Zhejiang University (Top 5%)

2020–2022

Second Prize, National College Mathematics Competition

2021